

SECTION 1 — MCP TOOLS (53 tools)

Category	Tool / Skill	Description
Core Routing	<u>route_request</u>	Route prompt to optimal model tier (Haiku/Sonnet/Opus) based on intent and stakes
	<u>analyze_tasks</u>	Route multiple tasks independently — each gets its own optimal model tier
	<u>rewrite_prompt</u>	Strip filler, tighten language, add role framing to reduce token waste
	<u>optimize_context</u>	Compress large context (>3K tokens) to fit a target token budget
	<u>plan_cache</u>	Design prompt-cache breakpoints for agentic or RAG repeated calls
	<u>batch_prompts</u>	Merge 2-5 <u>small related</u> tasks into one efficient batched request
	<u>summarize_thread</u>	Compress long conversation into a compact handoff summary for fresh threads
	<u>compare_providers</u>	Compare estimated cost across Claude, OpenAI, and Gemini tiers for a prompt
Session & Analytics	<u>ping_session</u>	Record session activity to reset idle clock and initialize cost tracking
	<u>check_session_timeout</u>	Check whether the current session is approaching its idle timeout
	<u>get_session_stats</u>	Display cumulative cost, savings, token burn rate, and role breakdown
	<u>get_routing_audit</u>	View routing decision log and why each model was chosen, plus pattern analysis
	<u>export_stats</u>	Export session usage history as JSON or CSV for offline analysis
	<u>clear_history</u>	Delete session history older than N days to free storage
	<u>reload_config</u>	Hot-reload pricing / providers / roles config without restarting the server
	<u>auto_compact</u>	Auto-compact threads approaching context window limits
Orchestration	<u>orchestrate_tasks</u>	Parse multi-task request, build dependency DAG, execute with optimal routing per node
	<u>auto_activate_plugins</u>	Detect which plugins apply to the current task and auto-load them
	<u>select_context_files</u>	Semantic search to pick the most relevant files to include in context
	<u>get_memory_snapshot</u>	Return current session memory state and key metrics in a single snapshot
	<u>export_session_summary</u>	Export full session summary (decisions, costs, models used) as JSON or CSV
	<u>view_dashboard</u>	Render live session dashboard and return its URL for browser viewing
	<u>compress_response</u>	Compress a text response at lite / full / ultra intensity levels
	<u>security_check</u>	Run a security validation pass on provided code before it ships
Skill System	<u>invoke_skill</u>	Load and execute a named <u>PromptWise</u> skill by ID
	<u>list_skills</u>	Discover available skills; filter by tag, role, or category
	<u>skill_chain</u>	Execute multiple skills sequentially, passing outputs as inputs
	<u>suggest_skill</u>	Recommend the best skill for a given user message or task description
Security	<u>check_prompt_injection</u>	Scan user input for prompt injection or jailbreak attempts before sending to model

	<u>suggest_skill</u>	Recommend the best skill for a given user message or task description
Security	<u>check_prompt_injection</u>	Scan user input for prompt injection or jailbreak attempts before sending to model
	<u>owasp_scan</u>	Scan code for OWASP Top-10 vulnerabilities and return ranked findings
	<u>scan_response</u>	Scan model response for PII leaks and injection echoes before surfacing to user
	<u>map_compliance</u>	Map security controls to SOC2 / NIST / ISO 27001 / GDPR frameworks
Prompt Registry	<u>register_prompt</u>	Register a new prompt or save a new version of an existing prompt to the registry
	<u>get_prompt</u>	Retrieve a registered prompt by name and optional version number
	<u>compare_prompts</u>	Diff two versions of a registered prompt and show token delta
AI Engineering	<u>audit_system_prompt</u>	Score system prompt on clarity, role, constraints, and jailbreak resistance
	<u>eval_prompt_across_models</u>	Estimate cost and recommended tier across Haiku / Sonnet / Opus
	<u>plan_context_window</u>	Plan optimal token allocation (system prompt, history, completion) given a budget
	<u>suggest_technique</u>	Auto-detect best prompt technique: CRAFT, Few-Shot, Chain-of-Thought, or Chaining
	<u>apply_craft</u>	Analyze prompt against CRAFT axes (Context/Role/Action/Format/Tone) and rebuild
	<u>inject_few_shot</u>	Enhance prompt with curated few-shot examples to anchor style or format
	<u>add_chain_of_thought</u>	Wrap prompt with Chain-of-Thought scaffolding to improve reasoning quality
	<u>chain_prompts</u>	Decompose complex task into a sequential prompt chain with defined handoffs
Dev Tools	<u>license_check</u>	Check OSS license compatibility for a list of packages against a target policy
	<u>cve_lookup</u>	Look up known CVEs for a package or dependency via the OSV.dev API
	<u>sbom_generate</u>	Generate a CycloneDX Software Bill of Materials for a project
	<u>generate_tests</u>	Generate a test skeleton (pytest / unittest / jest) from source code

	<u>analyze_coverage</u>	Parse <u>coverage.json</u> report and identify untested branches and gaps
	<u>github_action_check</u>	Scan repo for GitHub Actions workflows and verify <u>PromptWise</u> integration hooks
Budget	<u>predict_cost</u>	Estimate the cost of a prompt before sending it to the model
	<u>set_budget_limit</u>	Set monthly or daily spending limit in USD with configurable alert thresholds
	<u>get_budget_status</u>	Check current spend vs configured budget limits
	<u>budget_report</u>	Get detailed budget report with daily breakdown and trend analysis
Platform	<u>send_teams_alert</u>	Send a cost or threshold alert to a Microsoft Teams channel
	<u>send_slack_alert</u>	Send a cost or threshold alert to a Slack channel
	<u>trigger_github_workflow</u>	Trigger a GitHub Actions workflow programmatically from <u>PromptWise</u>
Enterprise	<u>create_team_session</u>	Create shared team session with role-based context and per-role budgets
	<u>get_team_stats</u>	Retrieve team-level usage statistics across all active members
	<u>set_role_budget</u>	Set per-role budget allocation (e.g., developer vs. architect quota)
	<u>export_team_report</u>	Export full team usage report for billing or audit purposes

SECTION 2 — SKILLS (30 skills)

Category	Tool / Skill	Description
Dev & Testing	<u>tdd</u>	Enforce Red → Green → Refactor TDD <u>cycle</u> ; failing test first, then minimal code
	test-generator	Generate unit and integration test skeletons from existing source code
	test-coverage-advisor	<u>Analyze</u> coverage gaps and generate targeted tests for untested branches
	feature-dev	Full feature lifecycle: brainstorm → TDD → code review → verify
	code-review	Structured review covering correctness, security, style, and test coverage
	refactoring	Guided refactoring with safety-net tests, incremental steps, and regression checks
Architecture	system-design	Generate C4 diagrams, data-flow maps, and deployment topology for a system
	architecture-review	Review for single points of failure, scalability issues, security gaps, and coupling
	<u>api-design</u>	Design REST or <u>GraphQL</u> APIs with versioning strategy, auth, and rate limiting
	<u>db-schema-design</u>	ERDs, normalization recommendations, indexing strategies, and partitioning guidance
	microservices-design	Define service boundaries, event contracts, saga patterns, and service mesh config
	agent-chain-designer	Design multi-agent workflows with coordination, handoff, and fallback logic
Data & AI	security-architecture	STRIDE threat <u>modeling</u> , define trust boundaries, and map mitigating controls
	data-flow-analysis	Map full data lifecycle including PII touchpoints, transformations, and retention
	rag-optimizer	Design or tune a RAG pipeline: chunk sizes, embeddings, reranking, retrieval quality
	few-shot-builder	Build and iteratively optimize few-shot example sets for a prompt
	system-prompt-auditor	Score a system prompt and return targeted rewrites for clarity and safety
	model-migration-advisor	Guide safe migration between Claude model versions with deprecation and cost analysis
Docs & Planning	<u>brd-generator</u>	Generate a Business Requirements Document from a product brief or stakeholder notes
	<u>prd-generator</u>	Generate PRD with goals, personas, success metrics, and scope
	user-story-generator	Generate user stories with acceptance criteria in Given / When / Then format
	<u>adr</u>	Author Architecture Decision Records capturing context, decision, and consequences
	<u>api-docs</u>	Generate or consume <u>OpenAPI</u> 3.0 specifications for a given API surface
	changelog-generator	Auto-generate a structured CHANGELOG from git log history
Process & Quality	runbook-generator	Generate ops runbooks and incident-response playbooks from a service description
	<u>systematic-debugging</u>	Structured debug flow: reproduce → isolate → hypothesize → fix → verify
	verification-before-completion	Pre-ship checklist: tests pass, security scan clean, docs updated, no regressions
	git-workflow	Enforce conventional commits, PR templates, and branch naming conventions
	meeting-notes	Structure a meeting transcript into decisions, action items, and owners